



QUIC
LEAGUE
2020

SAME SPORT, DIFFERENT STRATEGIES: **a study of QUIC and HTTP/3 game plan diversity**

ROBIN MARX, JORIS HERBOTS,
WIM LAMOTTE AND PETER QUAX

HASSELT UNIVERSITY - EDM

TCP
2U2
0P0

**TCP IS OUT,
QUIC IS IN!**



QUIC
LEAGUE
2020

**MORE OPTIONS FOR
INDIVIDUAL STACKS
TO TWEAK BEHAVIOUR**

- Stream multiplexing
- User space congestion control
- 0-RTT
- Binary framing

QUIC EVOLUTION

Google



VERSUS

Google

QUIC EVOLUTION



facebook

moz://a

Google



CLOUDFLARE®



fastly®

NetApp®



trafficserver™



And several others!

00

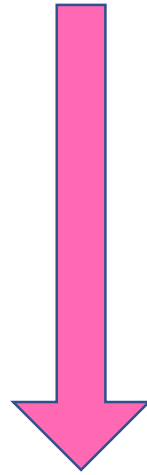
METHODOLOGY

ARE YOU WATCHING CLOSELY?

METHODOLOGY

structured
endpoint
logs

[qlog]



interactive
tooling

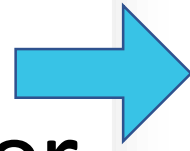
<qvis>

manual analysis

METHODOLOGY

QUIC
interop runner

- Automated tests, run daily
- Client-side behaviour

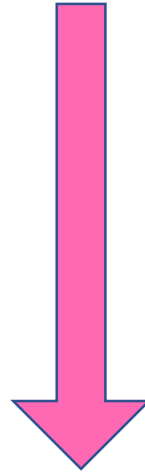


[qlog]



Custom
aioquic client

- Point at public interop endpoints
- Server-side behaviour



<qvis>

METHODOLOGY

QUIC interop runner

- Automated tests
run daily
- Client-side
behaviour

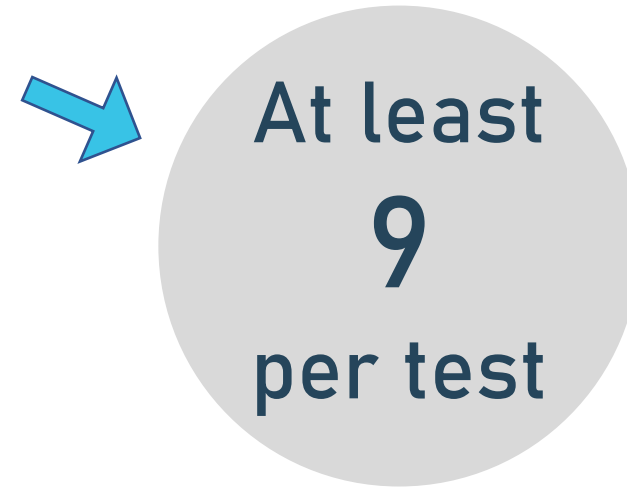
Verify results

1. Source code
review
2. Ask the
original
implementers

Custom quic client

point at public
interop endpoints
server-side
behaviour

METHODOLOGY



01

FLOW CONTROL

DON'T MAKE THAT BUFFER SUFFER

FLOW CONTROL

TCP
2 U 2
0 P 0

Single connection-level buffer
RECEIVE WINDOW



QUIC
LEAGUE
2020

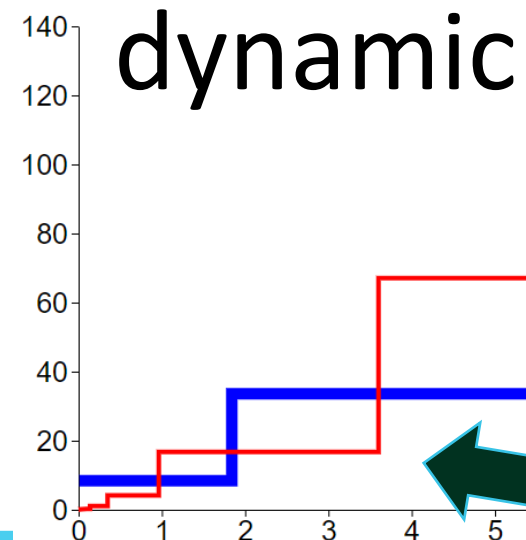
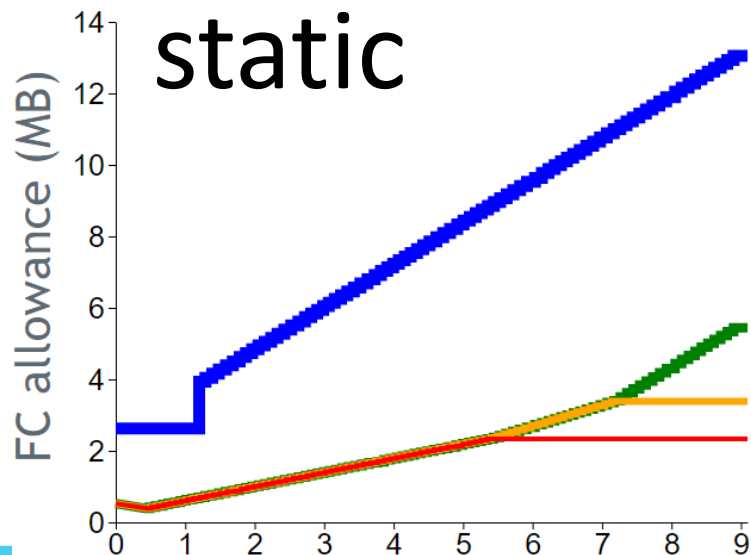
1. Connection-level limit
2. Stream-level limits
3. Stream count limit

} increase
often
OR
sending
can STALL

3 MAIN FC APPROACHES


1. static: 5000 received, you get 5000 more
2. dynamic: 5000 received, you get **10000** more
3. autotune: fluctuate based on RTT/application behavior

TCP
2U2
0P0

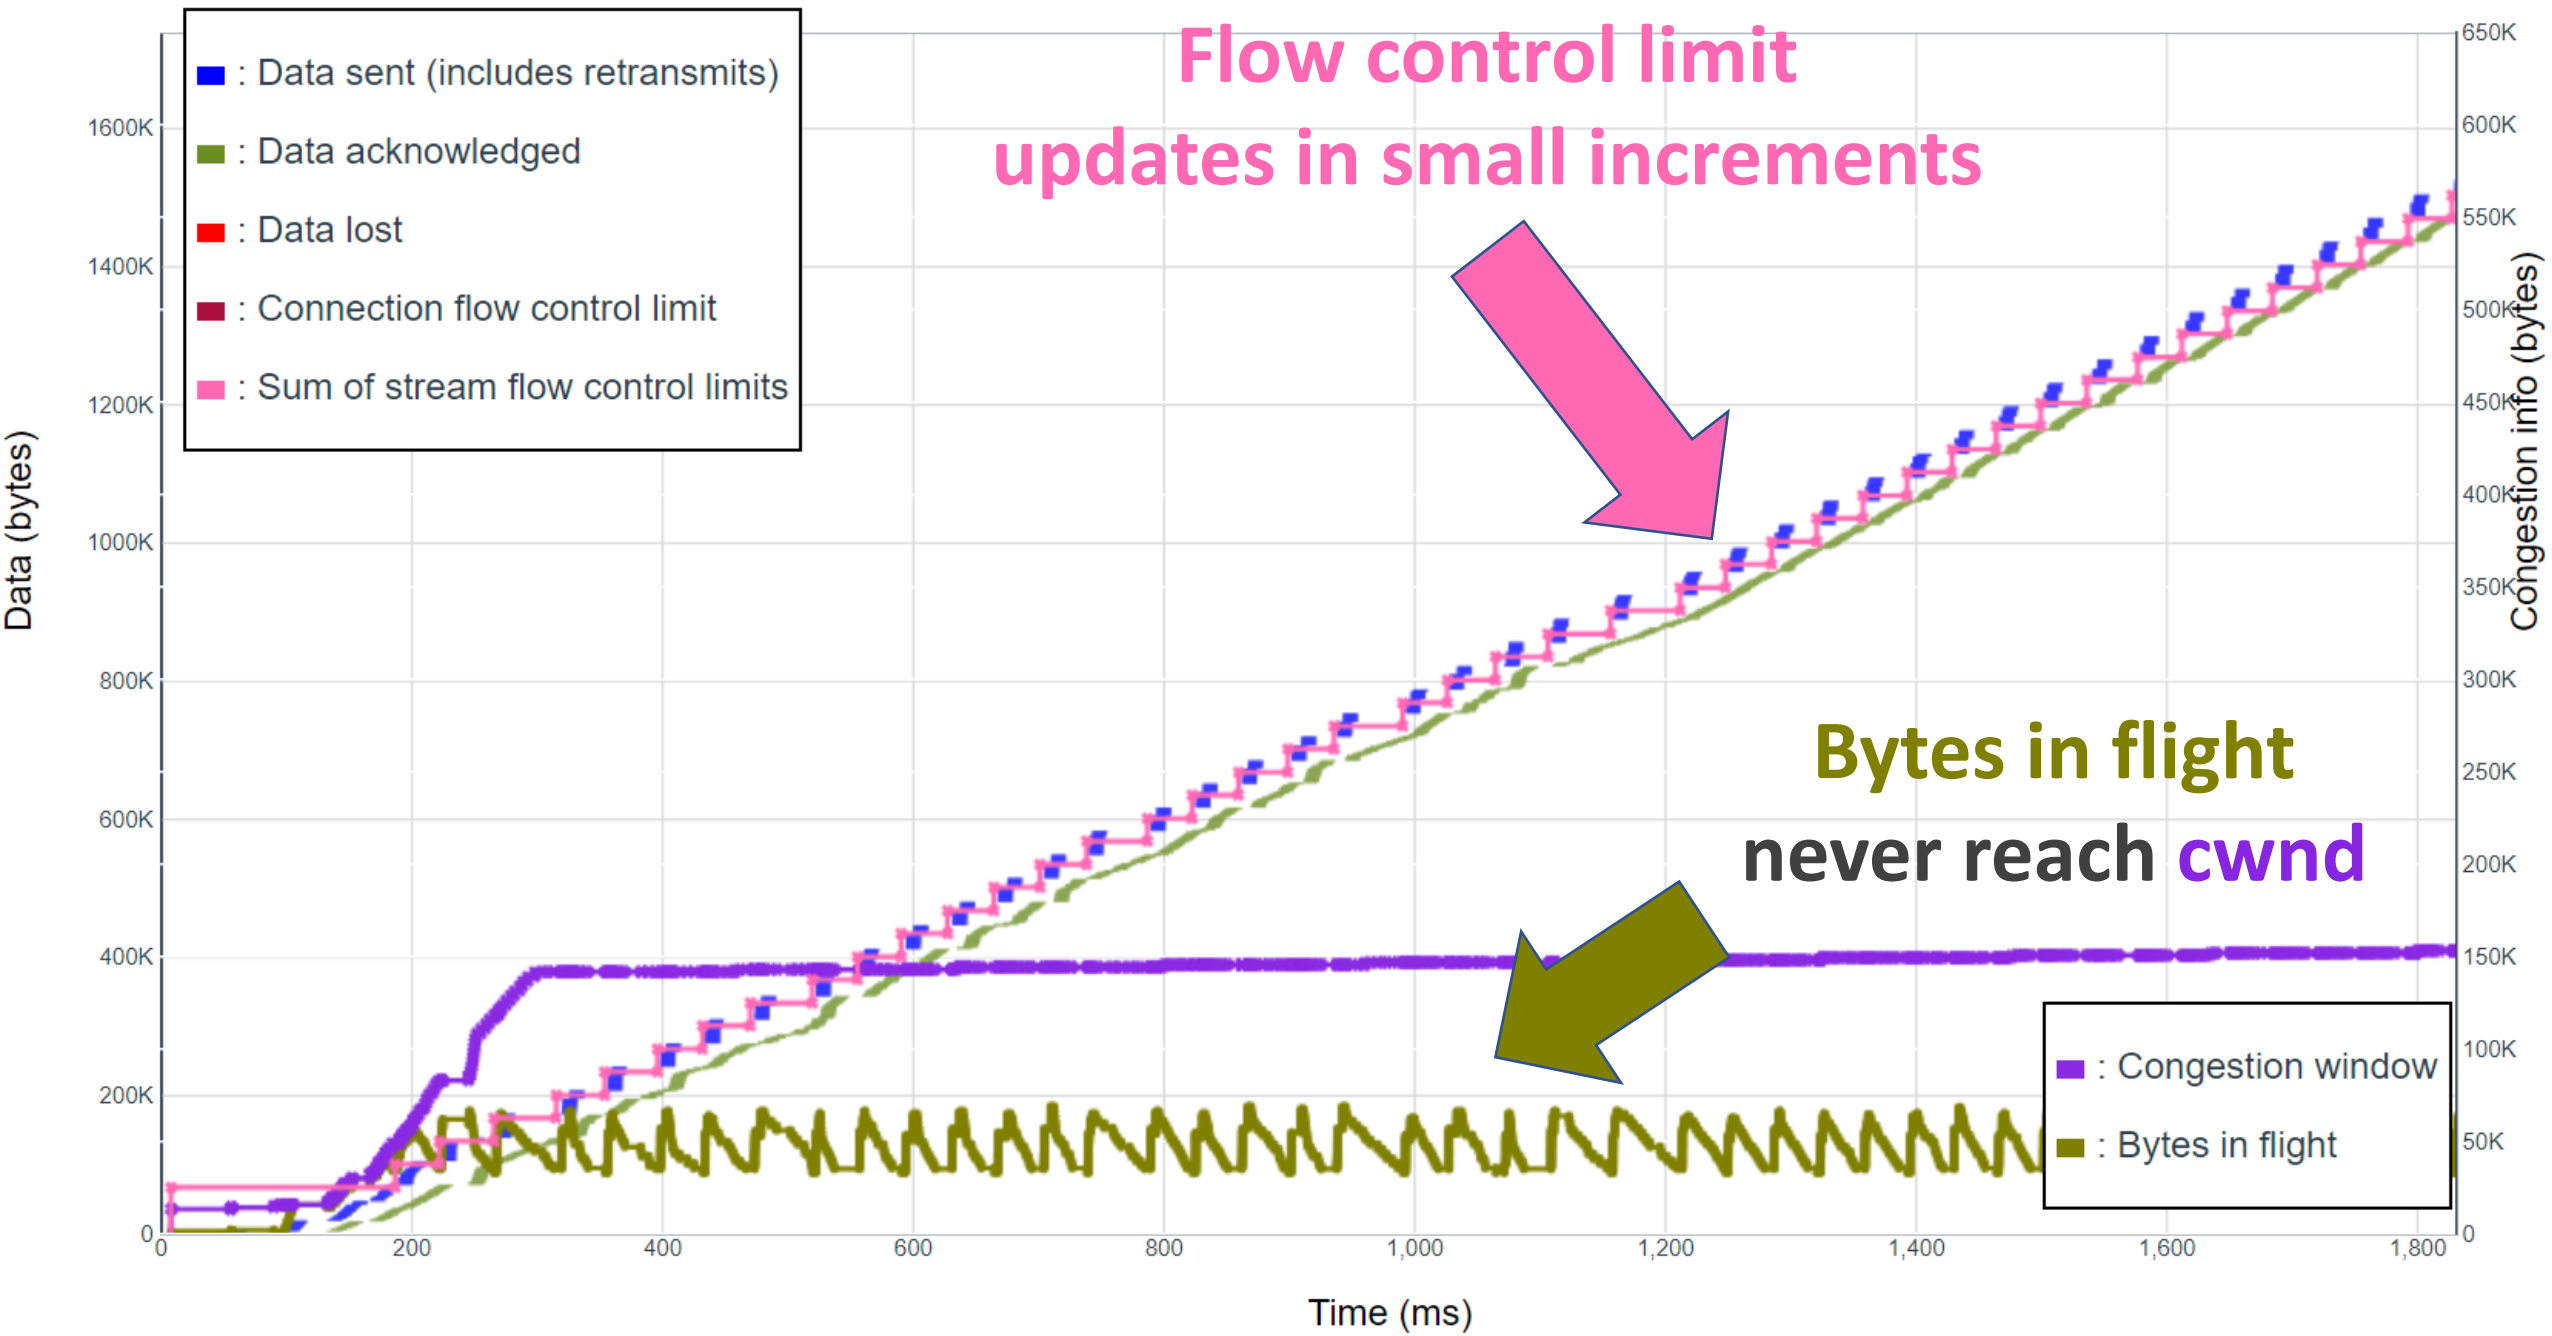


STALL

FLOW CONTROL

Flow control approach	Adoption	
static	8/12	
growing	3/12	
autotune	1/12	

“We have not yet spent time fine-tuning or testing Flow Control”

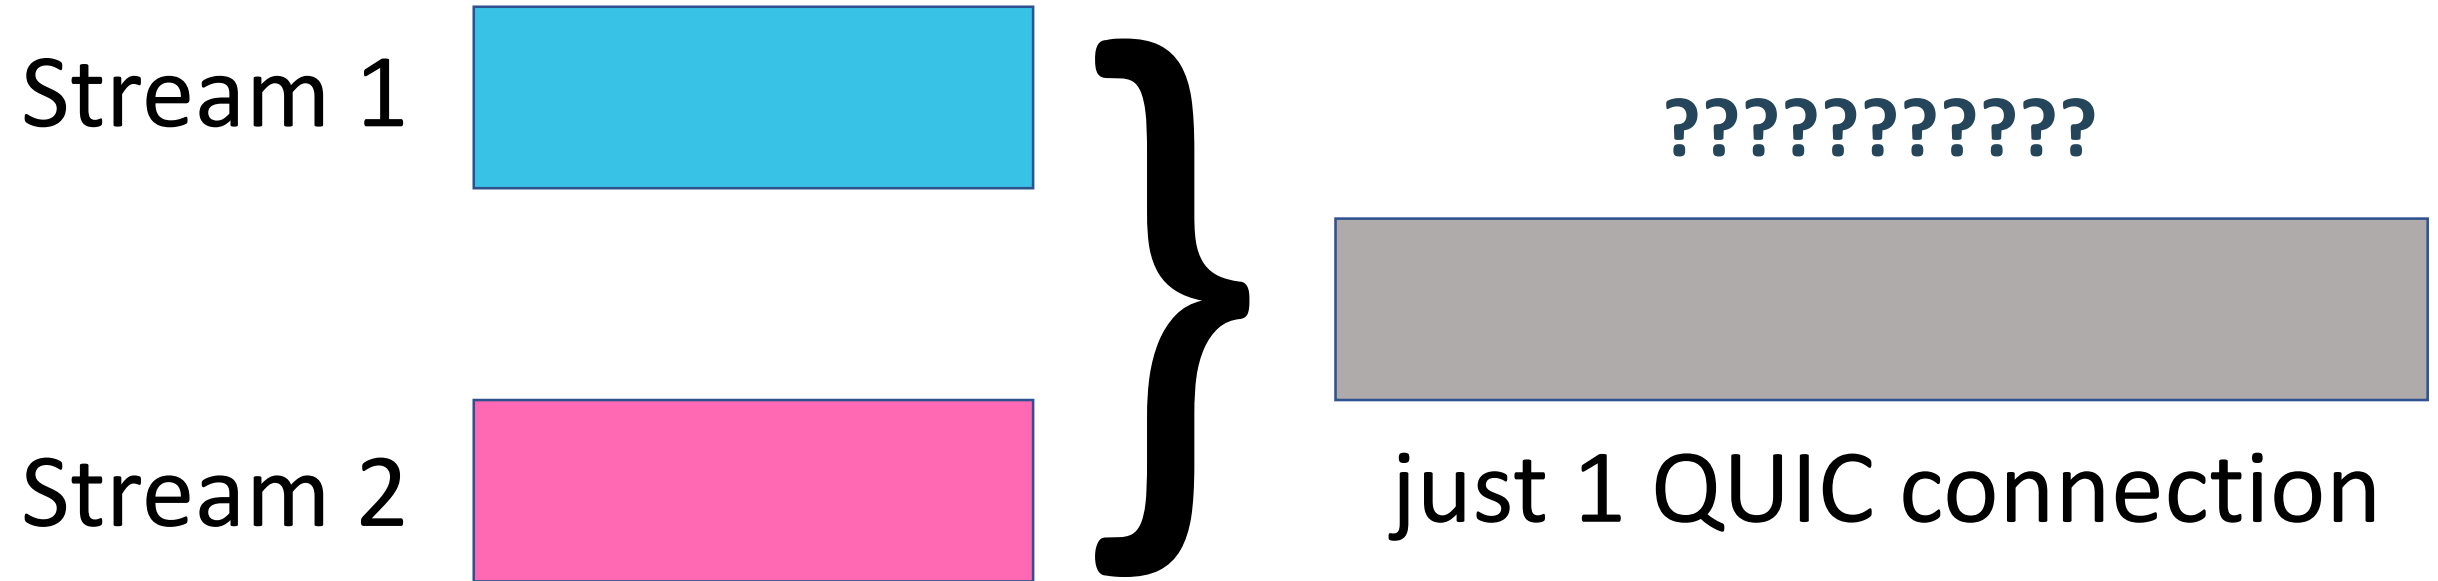


MULTIPLEXING

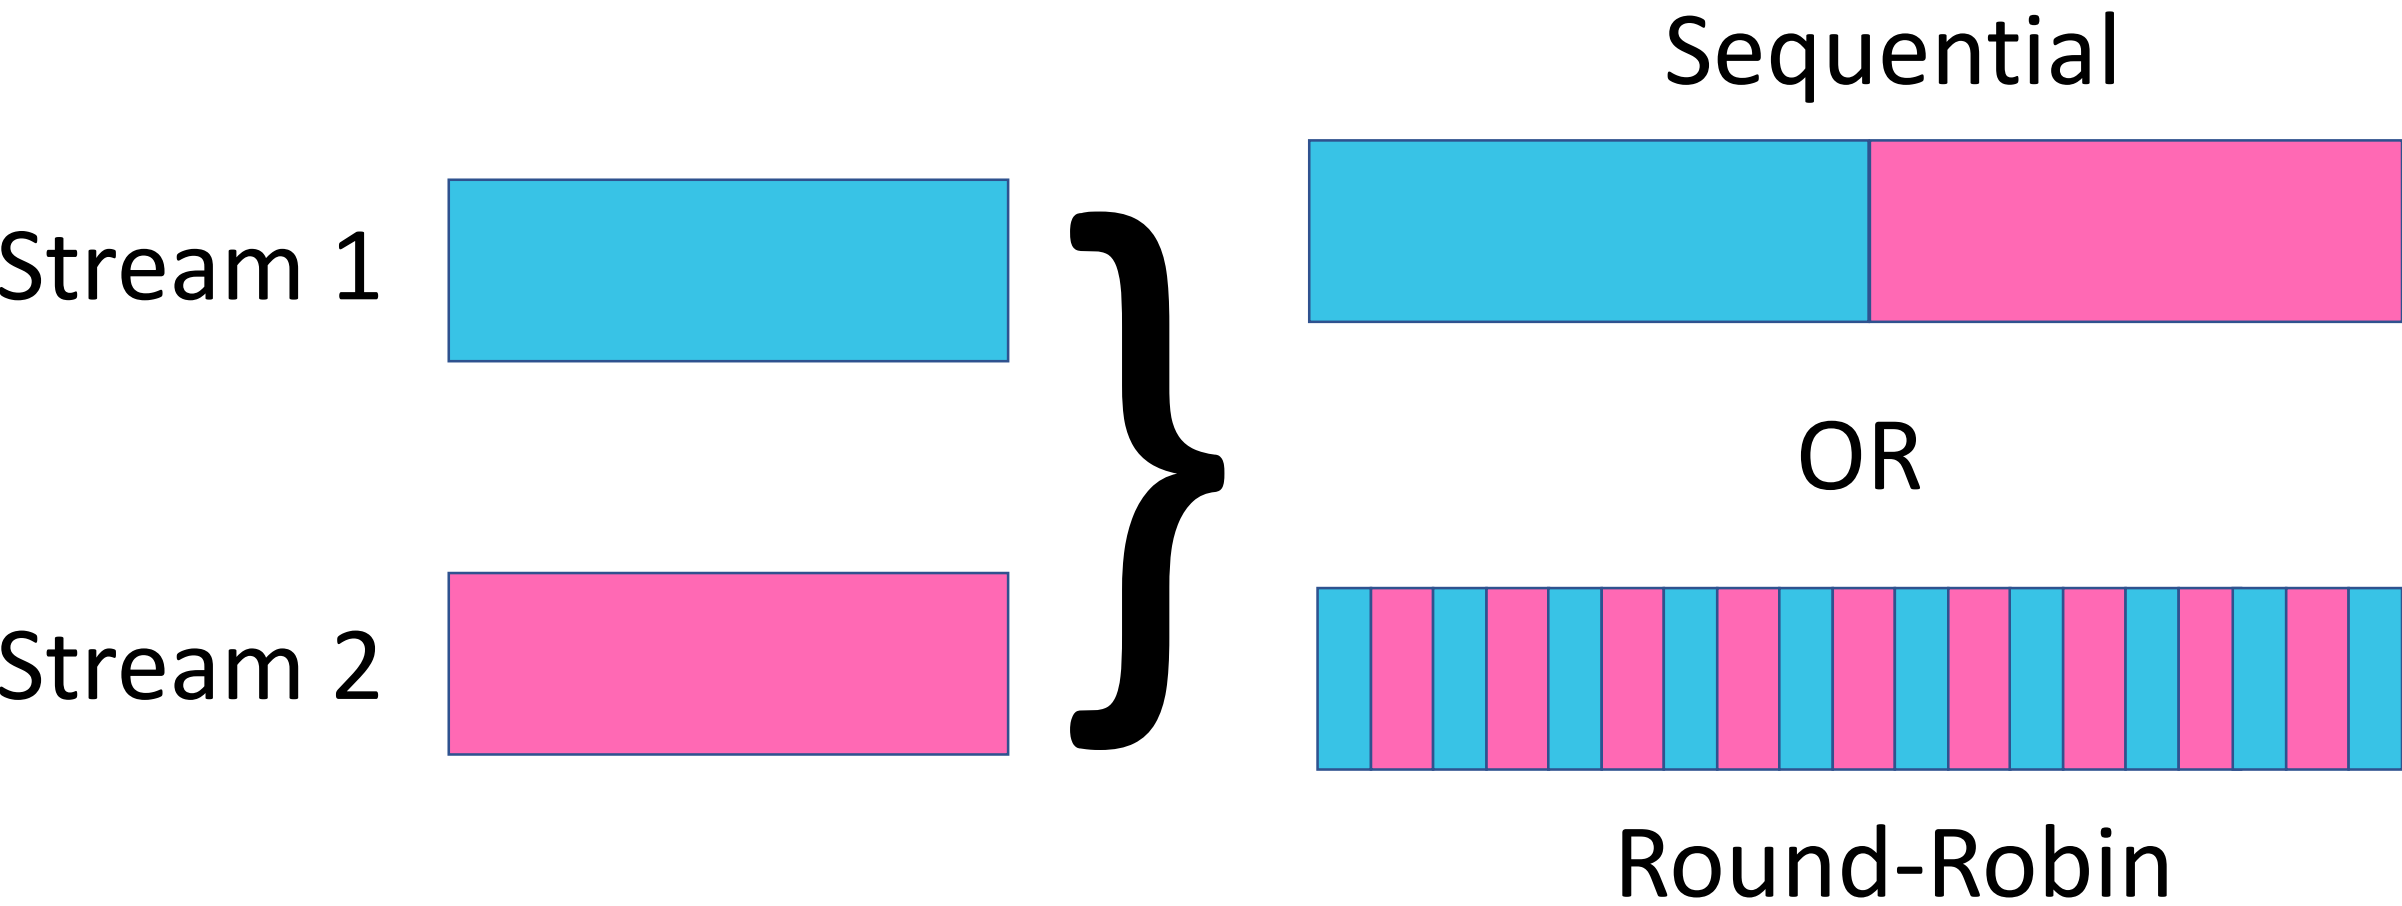
THE SPICE MUST FLOW

02

STREAM BANDWIDTH DISTRIBUTION

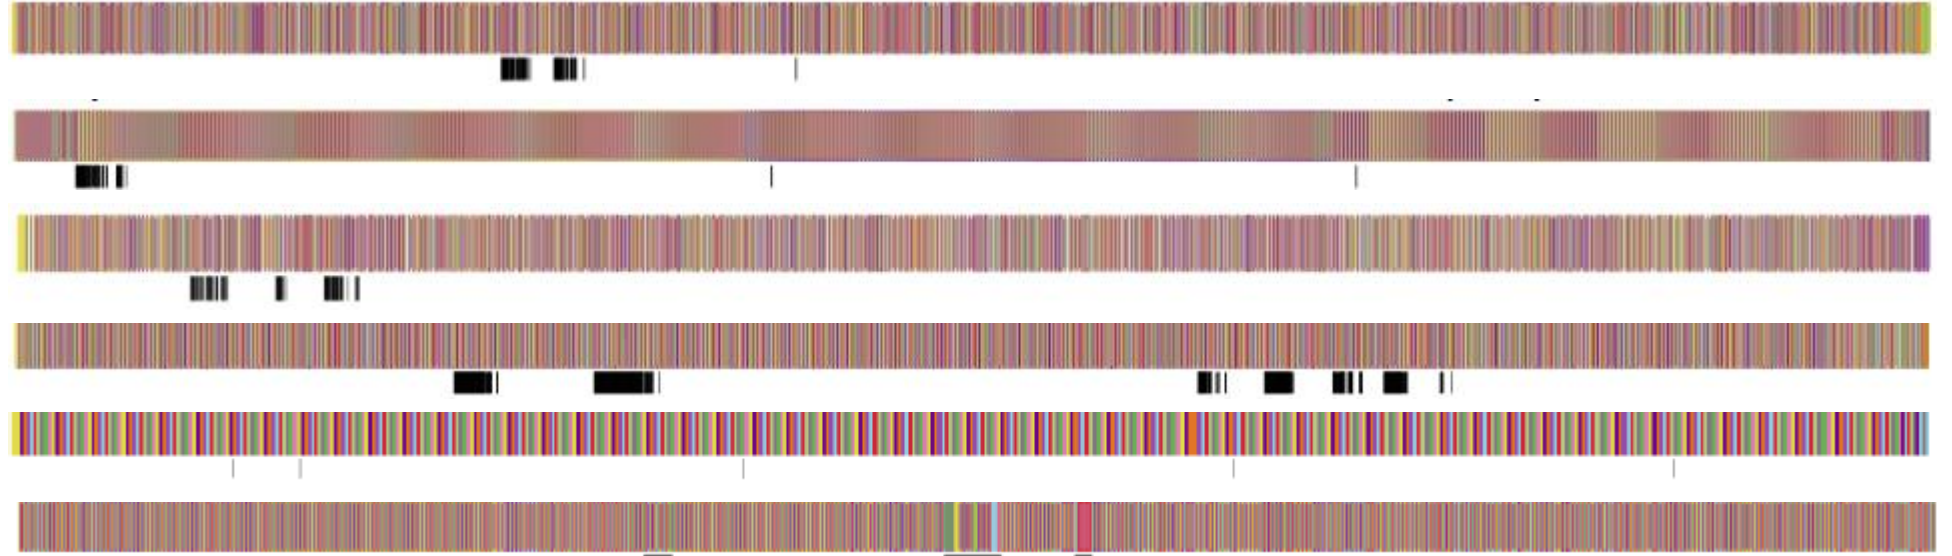


STREAM BANDWIDTH DISTRIBUTION

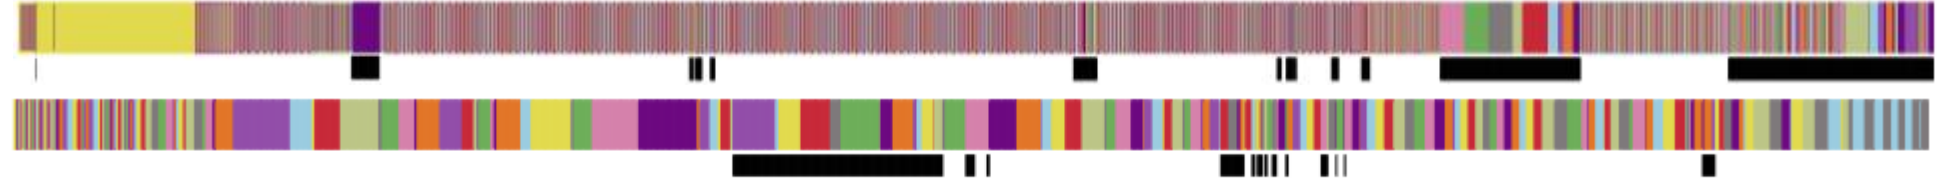


VARIETY IS THE SPICE OF LIFE

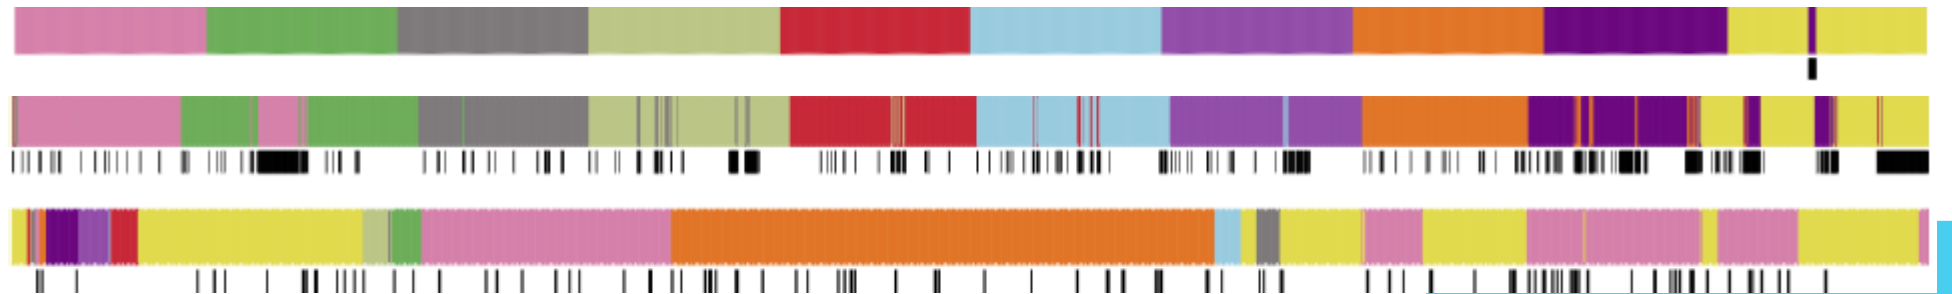
Round-Robin



The weird ones
in the middle



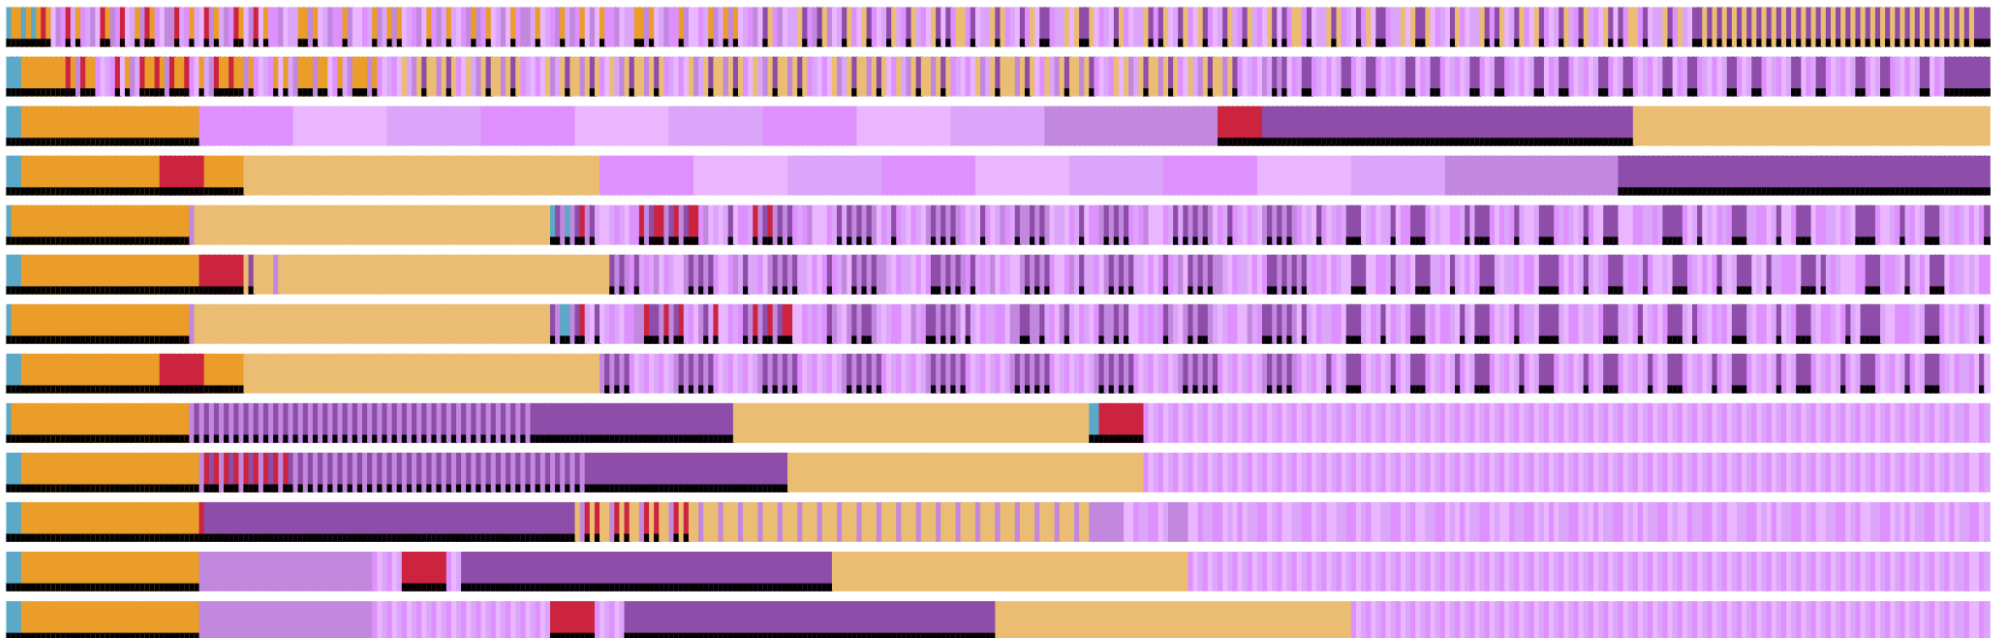
Sequential



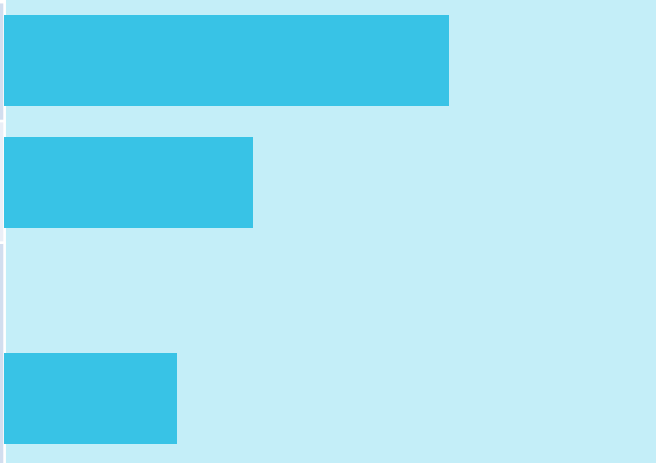

VARIETY IS THE SPICE OF LIFE



- Server Push
- Header Compression
- **Prioritization**



QUIC MULTIPLEXING

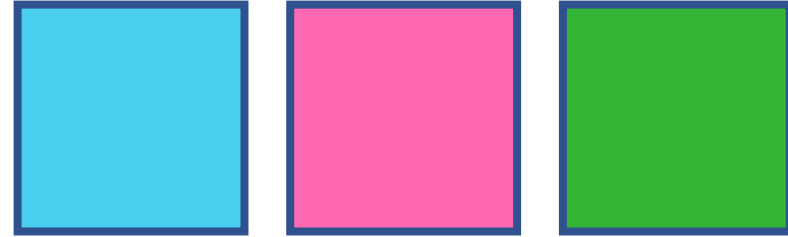
Multiplexer	Adoption	
Round-Robin	9/13	
Sequential	4/13	
 (experimental) HTTP/3 prioritization	5/18	

“waiting for HTTP/3 prioritization to fine-tune”

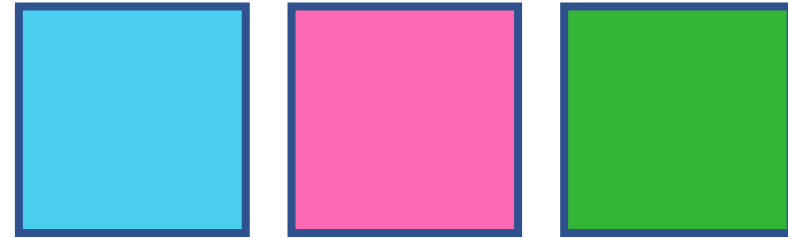
RETRANSMISSION SCHEDULING

TCP
2 U 2
0 P 0

send order



retransmission order



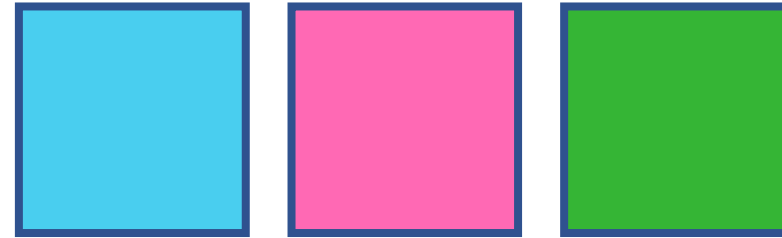
RETRANSMISSION SCHEDULING

TCP
2 U 2
0 P 0

send order



retransmission order

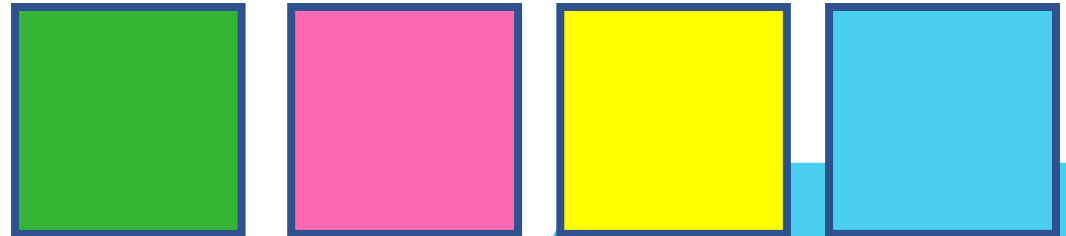


QUIC
LEAGUE
2020

send order



retransmission order



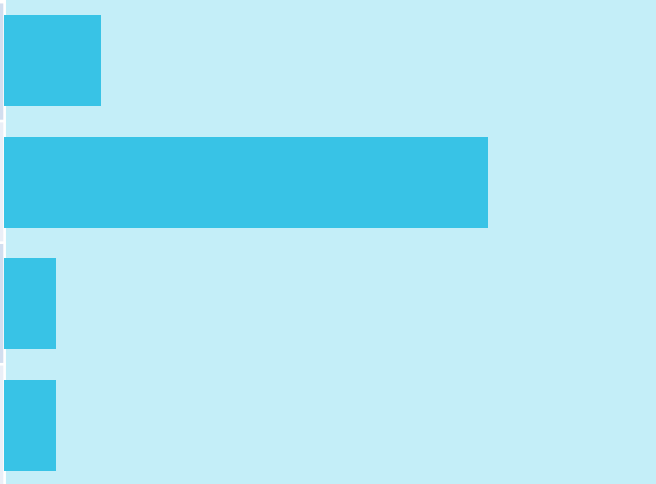

RETRANSMISSION SCHEDULING

1. No special treatment
2. Highest priority, default scheduler
3. Highest priority, different scheduler
4. (HTTP/3) Prioritization-driven

Example for nr. 3:



RETRANSMISSION SCHEDULING

Retransmission approach	Adoption	
1. All data is equal	2/13	
2. TCP-alike	9/13	
3. TCP-alike, change scheduler	1/13	
4. Prioritization-driven 	1/13	

“Unclear which performs best/if it matters”

03

CONGESTION CONTROL

THE NEED FOR SPEED

CONGESTION CONTROL

(MOST INACCURATE TIMELINE EVER)

TCP
2U2
0P0

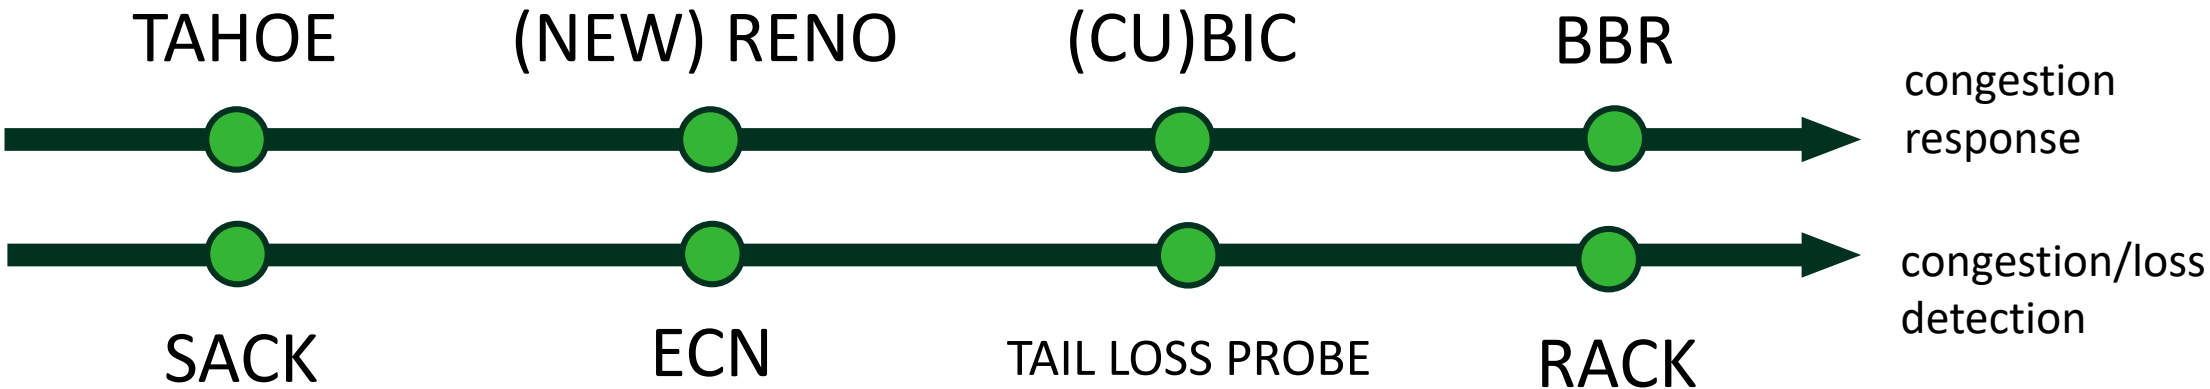


QUIC
LEAGUE
2020



CONGESTION CONTROL (MOST INACCURATE TIMELINE EVER)


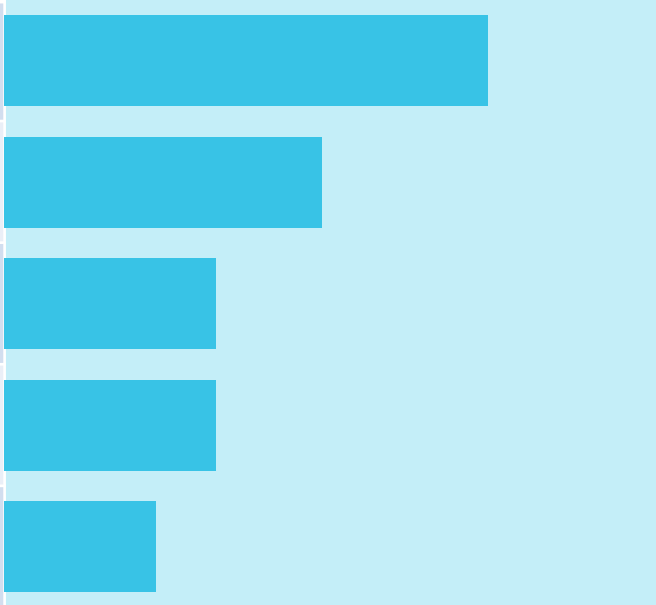
TCP
2U2
0P0



QUIC
LEAGUE
2020



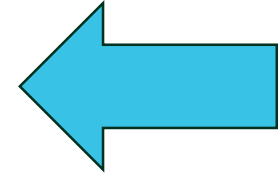
CONGESTION CONTROL

Congestion controller	Adoption	
New Reno 	9/15	
CUBIC	6/15	
→ (with hystart)	→ 4/6	
BBR v1	4/15	
BBR v2, COPA, ...	3/15	

“Often too complex to implement a new one”



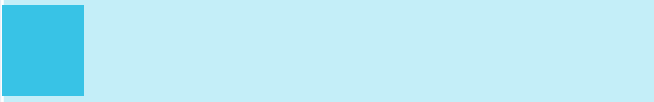
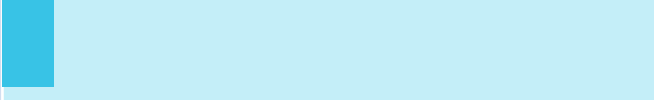
CONGESTION CONTROL

```
congestion := congestion.NewCubicSender(  
    congestion.DefaultClock{},  
    rttStats,  
    true, // use Reno  
    tracer,  
)
```






SNEAKY SNEAKY

THE DEVIL IS IN THE DETAILS


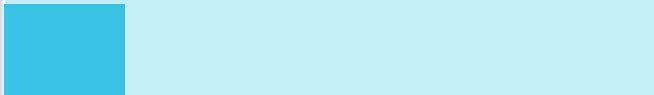

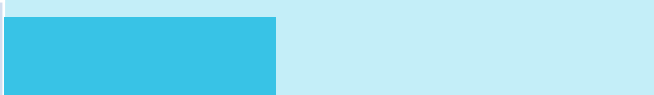
Initial congestion window	Adoption	
12-15 kB 	11/14	
40+ kB	3/14	
smart tweaking	2/14	

“We just looked at what Google was doing”

Pacing	Adoption	
Yes 	8/15	
No	7/15	

“Complex to get right”

THE DEVIL IS IN THE DETAILS

ACK every X packets	Adoption	
2 	2/12	
1 - 38	10/12	
ACK frequency extension	4/12	

“Read from socket in large batches, ACK per batch”

“Lower ACK frequencies are better on constrained networks”

0-RTT

—
OFF-BY-ONE ERRORS ARE THE BEST

04

ROUND TRIPS ARE THE WORST

TCP
2 U 2
0 P 0

1. SYN/ACK
2. TLS
3. (TLS)
4. HTTP



1. QUIC + TLS
2. HTTP

1. QUIC + TLS + HTTP

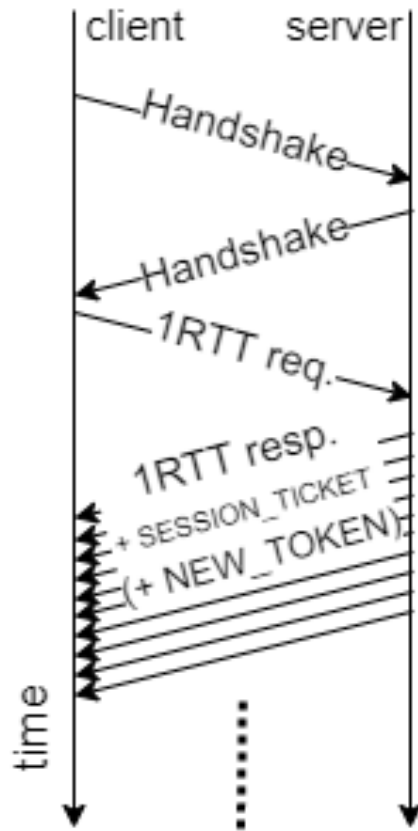
SOCCER HAS OFFSIDE, QUIC HAS 0-RTT

QUIC 0-RTT:

1. Needs to be encrypted
 - Only from second connection (session ticket)
2. Runs over IP + UDP
 - Send max **3X** as much as received (**amplification limit**)
3. Transports HTTP
 - Only idempotent requests

MAKING 0-RTT BETTER

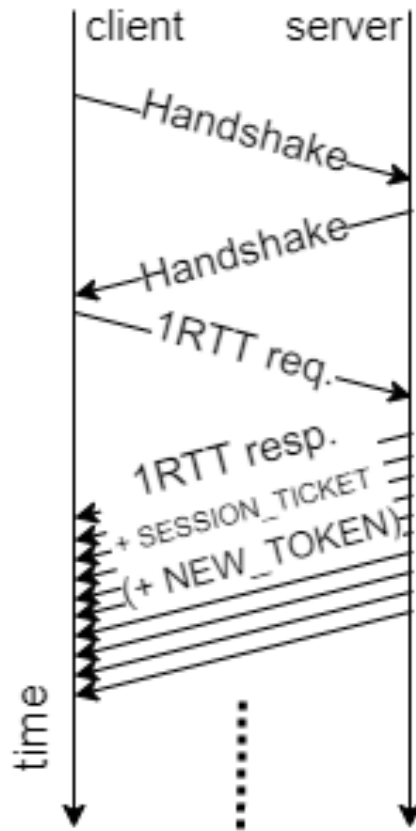
1-RTT



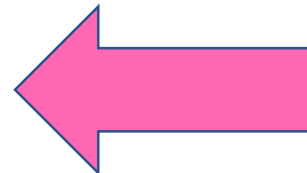
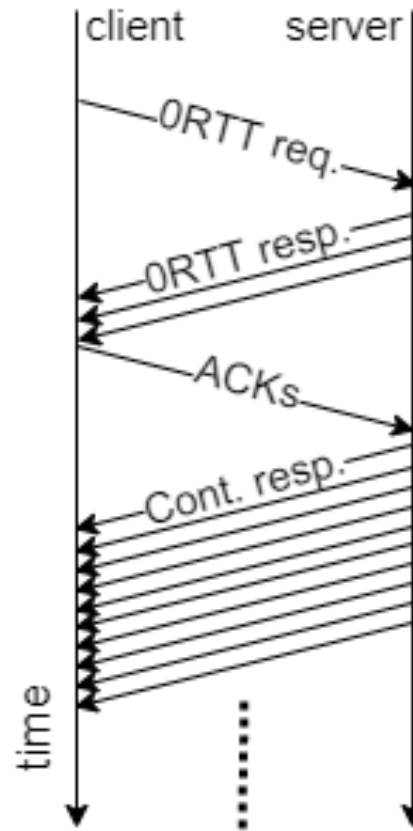
Session ticket enables 0-RTT
for the **next** connection

MAKING 0-RTT BETTER

1-RTT



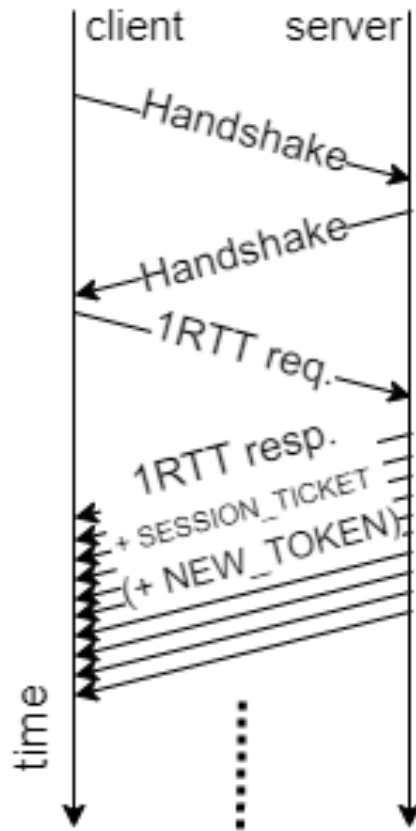
0-RTT



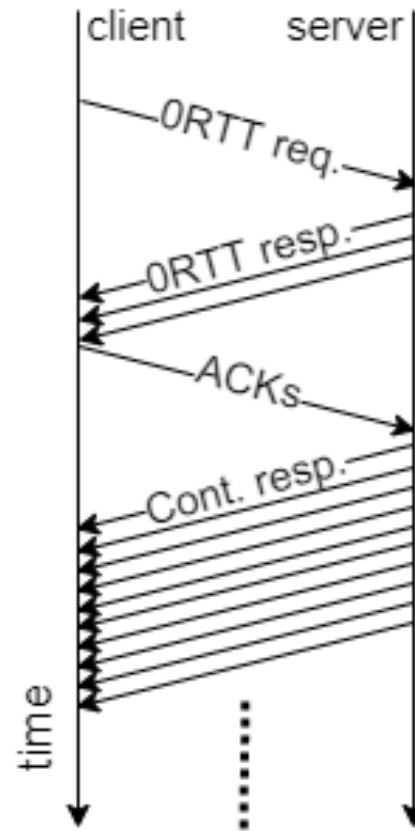
Server can only send **3X** what it received (say 4-5 kB)

MAKING 0-RTT BETTER

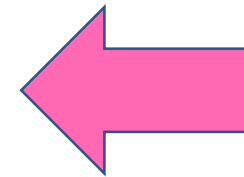
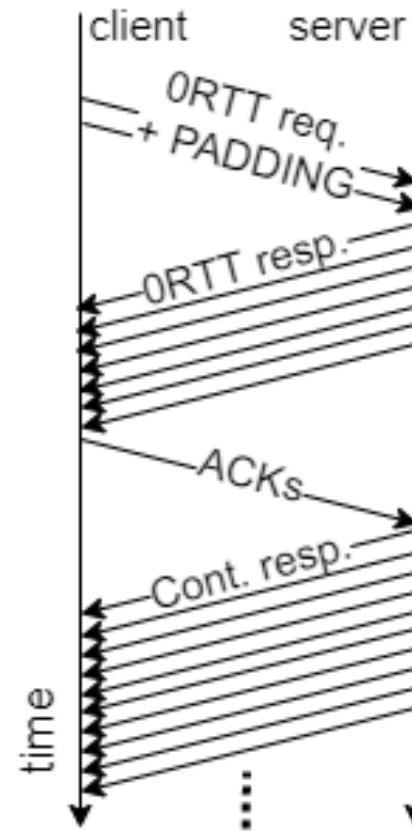
1-RTT



0-RTT



0-RTT with padding

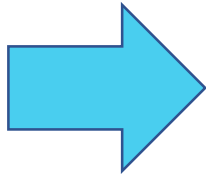


Client sends
more, server
can send
more

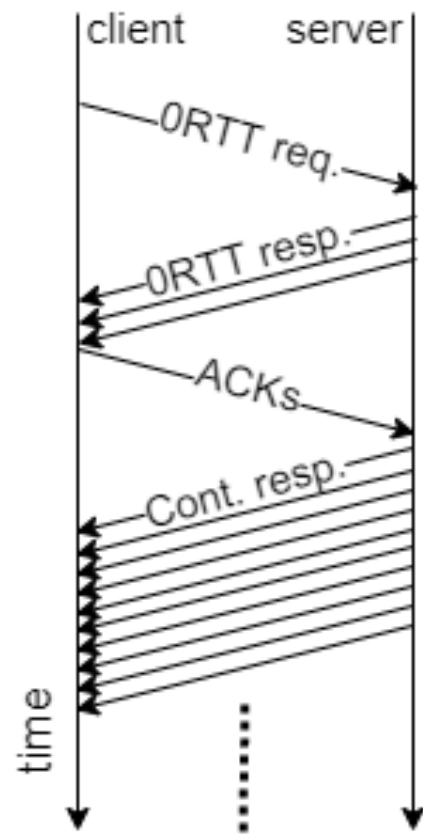
MAKING 0-RTT BETTER

1-RTT

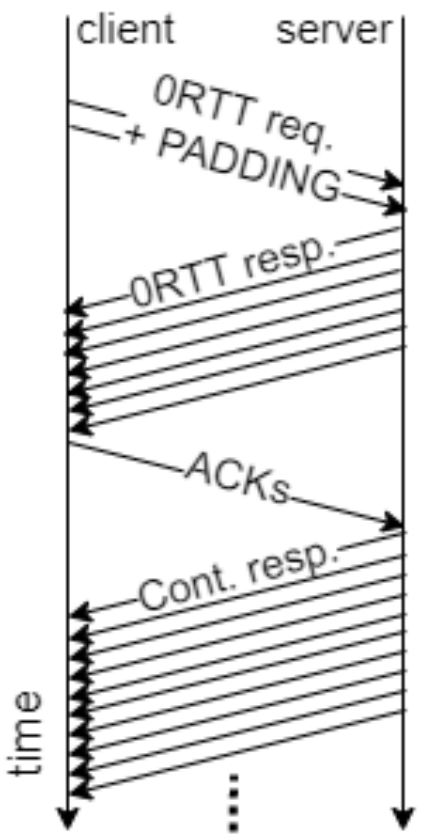
Address
validation
token



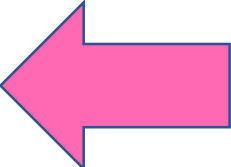
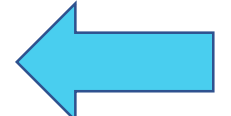
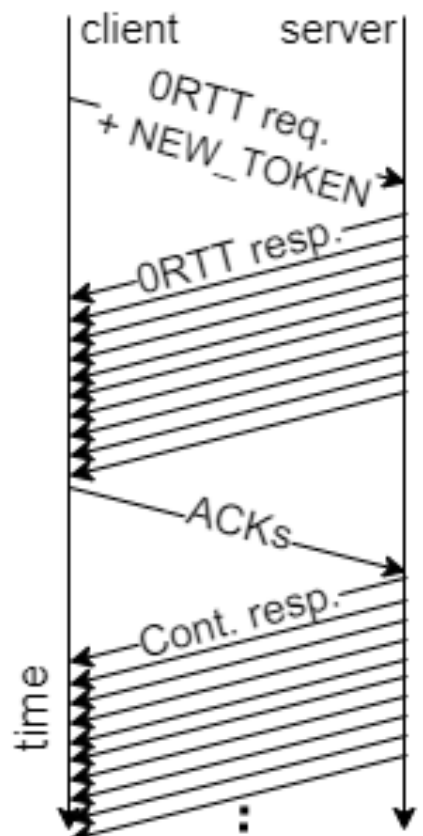
0-RTT



0-RTT
with padding




0-RTT
with address token




Send up
to initial
cwnd

THE OPTIMIZATION FORMERLY KNOWN AS CRUCIAL

0-RTT support	Adoption	
Yes 	13/18	<div><div></div></div>
No	5/18	<div><div></div></div>

“TLS library doesn’t support it yet”

Optimizations	Adoption	
Extra PADDING	0/9	<div><div></div></div>
NEW_TOKEN 	7/13	<div><div></div></div>

Amplification bugs	4/9	<div><div></div></div>
--------------------	-----	------------------------

SET THE LASERS TO AMPLIFICATION

1. Ignore limit, have a 46kB init cwnd
= **36X** amplification
2. Do not apply limit to retransmissions of 0-RTT data
= **17X** amplification
3. Do not apply congestion control to 0-RTT data
= **300kB burst if client sends 100kB**



05

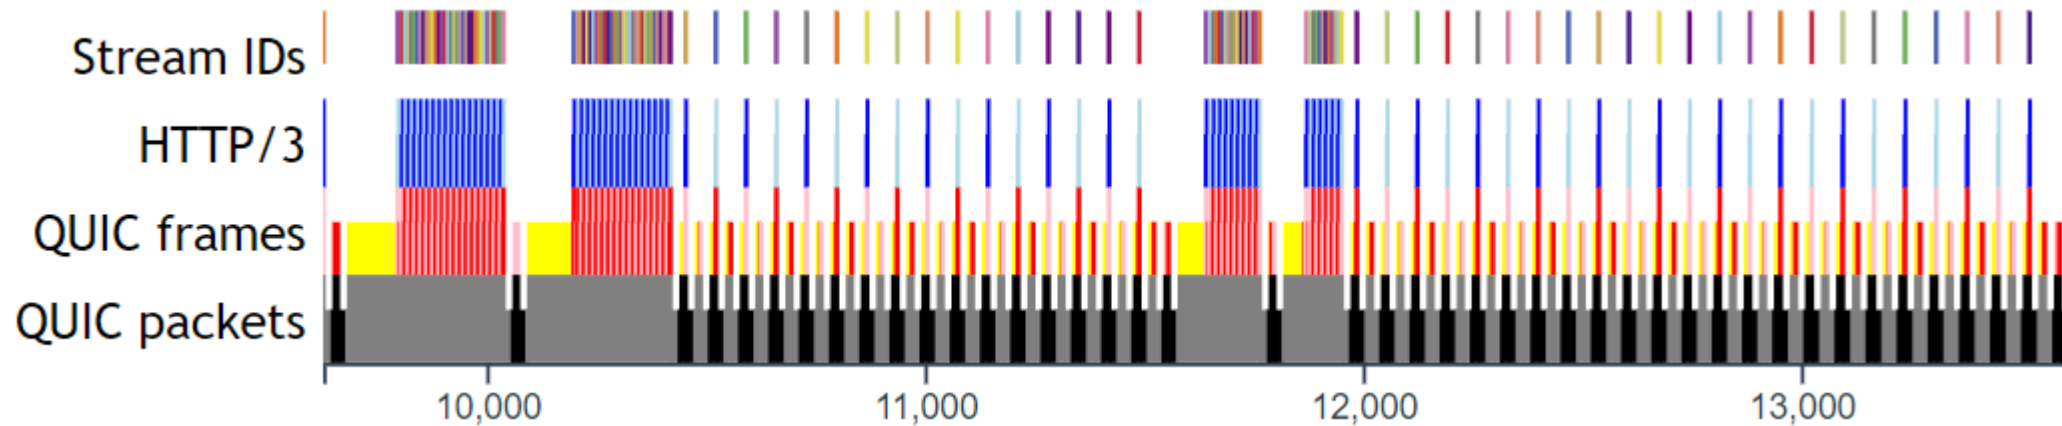
PACKETIZATION EFFICIENCY

THINKING OUTSIDE OF THE BOX

FRAMECEPTION

HTTP/3 **frames** are packet inside
QUIC **frames** which are packet inside
QUIC **packets**, which are sized after **MTU discovery**

→ **Efficiency + Head-of-Line blocking**





QUIC
LEAGUE
2020

CONCLUSION

IT'S JUST A GAME

LOOK BEFORE YOU LEAP

1. QUIC is complex:

many knobs to turn, easy to make it slower/faster

2. QUIC and HTTP/3 implementations aren't finished:

don't trust, always verify

QUIC 1 != QUIC 2 != QUIC 3 != QUIC 4

test different implementations

3. You might want to look at our methodology

[qlog] <qvis>

IMAGE SOURCES

- <https://seeklogo.com/images/V/versus-logo-97EAA46E88-seeklogo.com.png>

CREDITS & COPYRIGHTS

- Template created by Showeet.com
- **Free font used:**
 - Calibri (Microsoft font)
- **Copyrights:**
 - Free with Attribution
 - Cannot be resold or redistributed under any circumstances
 - Cf. <https://www.showeet.com/terms-of-use/>

THANK **Y**OU!

Hope you like this template :)



Free creative templates,
charts, diagrams and maps
for your outstanding
presentations

