


SmartEntry: Mitigating Routing Update Overhead with Reinforcement Learning for Traffic Engineering

Junjie Zhang, Zehua Guo, Minghao Ye, H. Jonathan Chao




Background

- **Traffic Engineering (TE):** Configure routing to improve network performance
- **Metric:** Maximum Link Utilization (MLU) $\rightarrow \frac{\text{Load}}{\text{Capacity}}$ of the most congested link




Background

- **Traffic Engineering (TE):** Configure routing to improve network performance
- **Metric:** Maximum Link Utilization (MLU) $\rightarrow \frac{\text{Load}}{\text{Capacity}}$ of the most congested link




Background

- **Traffic Engineering (TE):** Configure routing to improve network performance
- **Metric:** Maximum Link Utilization (MLU) $\rightarrow \frac{\text{Load}}{\text{Capacity}}$ of the most congested link



Flow-based or Destination-based Routing?

Flow-based Routing



Flow table at node 5


Match	Action
src = 2, dst = 7	Fwd to 6
src = 4, dst = 7	Fwd to 8

Two different sources can reach the same destination with different preconfigured paths

- 😊 Fine-grained traffic distribution control
- ☹️ Need to store $O(P^2)$ flow entries!
Scalability issue with limited TCAM resource

$P = \# \text{ of IP routes}$

Destination-based Routing



Forwarding table at node 5

Destination	Next Hop
7	6

Paths from two different sources to same destination must coincide once they overlap

- 😊 ➤ Lower forwarding complexity - $O(P)$ entries
- Widely implemented with simple RAMs

Centralized controller can be applied to update the entries when traffic changes

Motivation

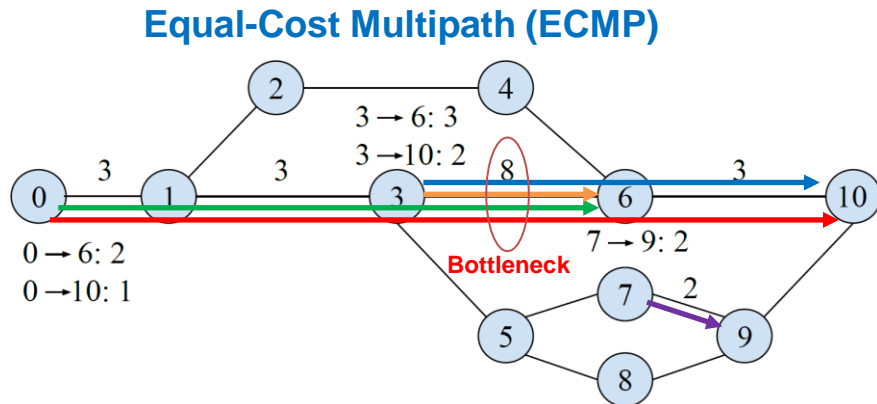
However, traditional TE need to update all entries to improve network performance!

☹ Take considerable time → cannot react to traffic changes in a responsive manner

Q: Can we mitigate routing update overhead?

A: Differentiate and route flows with a new traffic abstraction!

- (1) Only update some **critical entries** at some **critical nodes** to reroute traffic
- (2) The remaining unaffected traffic are forwarded by ECMP



Motivation

However, traditional TE need to update all entries to improve network performance!

☹ Take considerable time → cannot react to traffic changes in a responsive manner


Q: Can we mitigate routing update overhead?

A: Differentiate and route flows with a new traffic abstraction!

- (1) Only update some **critical entries** at some **critical nodes** to reroute traffic
- (2) The remaining unaffected traffic are forwarded by ECMP

Forwarding table at node 1

Destination	Next Hop
6	2 (100%)



Motivation

However, traditional TE need to update all entries to improve network performance!


☹️ Take considerable time → cannot react to traffic changes in a responsive manner

Q: Can we mitigate routing update overhead?

**Key Problem: which pairs are 'critical'?
There are too many (node, dst) combinations!**

A: Differentiate and route flows with a new traffic abstraction!

- (1) Only update some **critical entries** at some **critical nodes** to reroute traffic
- (2) The remaining unaffected traffic are forwarded by ECMP



Forwarding table at node 1

Destination	Next Hop
6	2 (100%)

Forwarding table at node 3

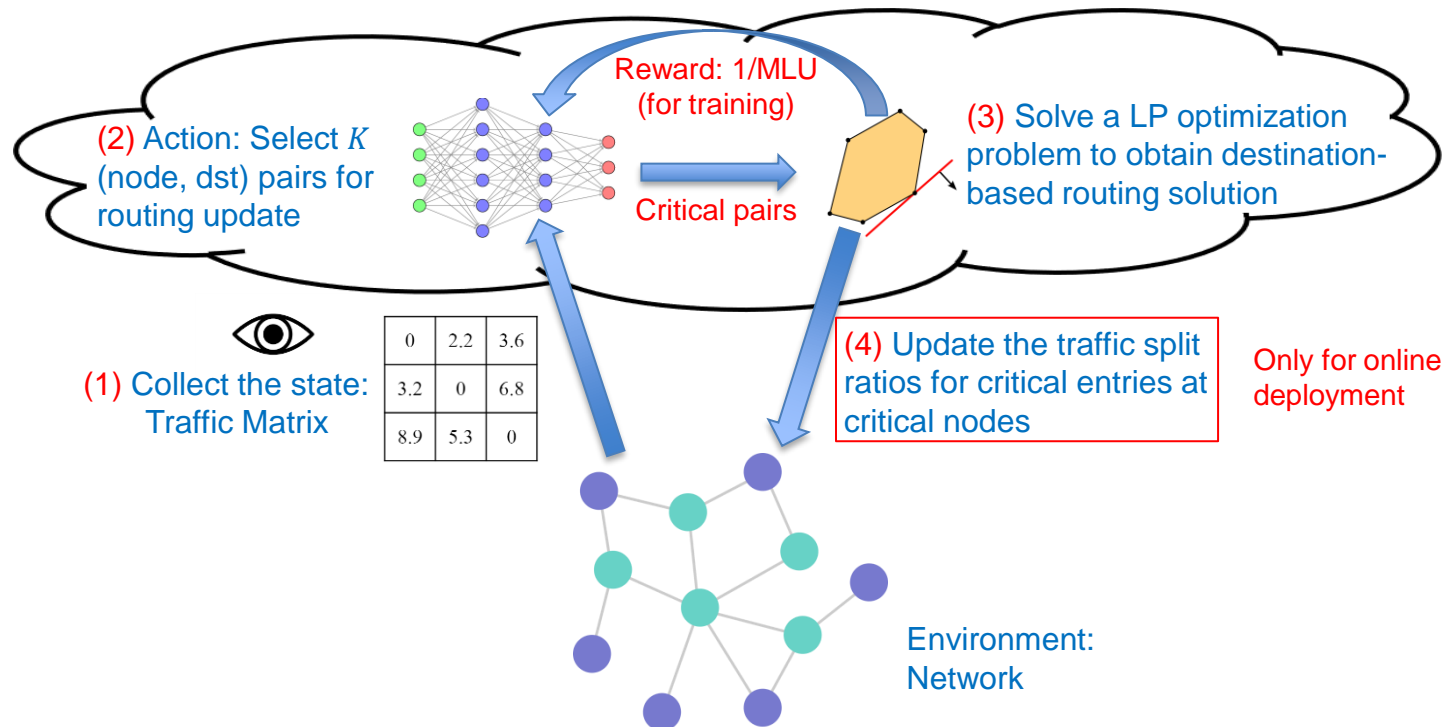
Destination	Next Hop
10	5 (100%)

Forwarding table at node 5

Destination	Next Hop
10	7 (33.3%), 8 (66.6%)

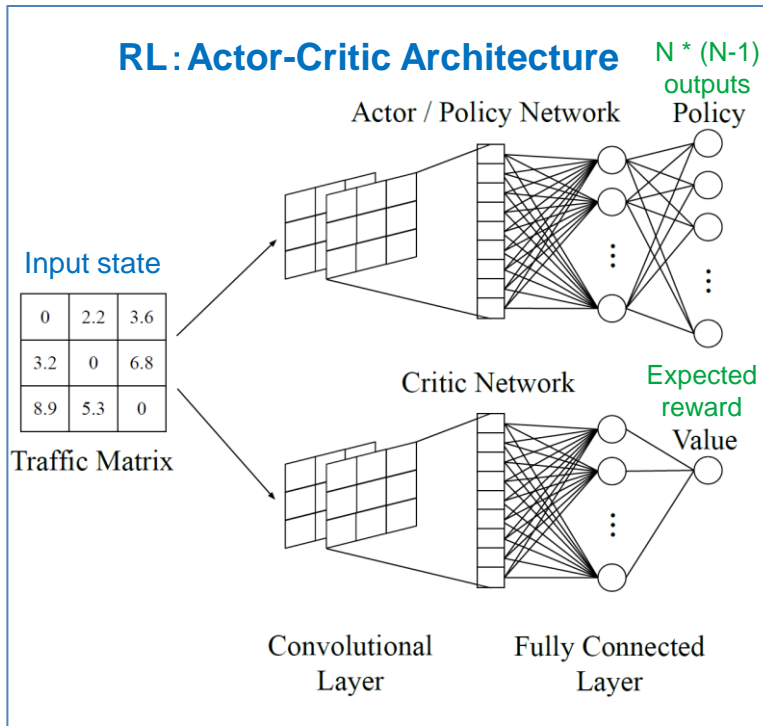
SmartEntry: RL + LP combined approach

- Idea: (1) Using **Reinforcement Learning (RL)** to smartly select critical pairs for routing update
(2) Solve a **Linear Programming (LP)** optimization problem to obtain destination-based routing solution



Why is RL + LP powerful?

- RL can model complex selection policies as neural networks to map “raw” observations to actions
- LP generates reward signal for RL to learn a better combination selection policy (minimize MLU)



Gradient update

Actions

LP minimize $U_{max} + \epsilon \cdot \sum_{\langle i,j \rangle \in E} \sum_{d \in V} y_{i,j}^d$ Produce reward signal

subject to

$$\sum_{d \in V} y_{i,j}^d = l_{i,j} \quad i, j : \langle i, j \rangle \in E$$

$$l_{i,j} \leq c_{i,j} \cdot U_{max} \quad i, j : \langle i, j \rangle \in E$$

$$\sum_{k: \langle k, i \rangle \in E} y_{k,i}^d - \sum_{k: \langle i, k \rangle \in E} y_{i,k}^d = -t^{i,d} \quad i, d : \tau_i^d \in \tau_K$$

$$y_{i,k}^d = \begin{cases} \frac{\sum_{n: \langle n, i \rangle \in E} y_{n,i}^d + t^{i,d}}{|\text{ENH}_i^d|} & \text{if } k \in \text{ENH}_i^d \\ 0 & \text{otherwise} \end{cases}$$

$$i, d : \tau_i^d \in \tau_{N \times (N-1) - K}, k : \langle i, k \rangle \in E$$

$$\sum_{k: \langle k, d \rangle \in E} y_{k,d}^d - \sum_{k: \langle d, k \rangle \in E} y_{d,k}^d = \sum_{s \in V, s \neq d} t^{s,d} \quad d \in V$$

$$y_{i,j}^d \geq 0 \quad d \in V, i, j : \langle i, j \rangle \in E$$

6

Experiment setup

➤ We use 4 real networks to evaluate SmartEntry



➤ Baseline Methods

- ❖ ECMP: Distributes traffic evenly among available next hops along the shortest paths
- ❖ Weighted ECMP: extends ECMP to allow weighted traffic splitting with shortest paths

➤ Evaluation Metric: **Performance Ratio (PR)**

- ❖ Compare against optimal flow-based routing in terms of MLU
- ❖ $PR = MLU_{optimal} / MLU_{SmartEntry}$

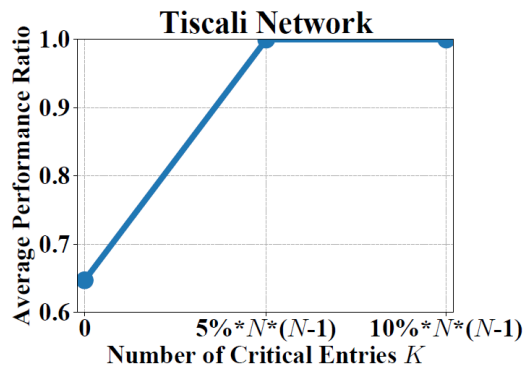
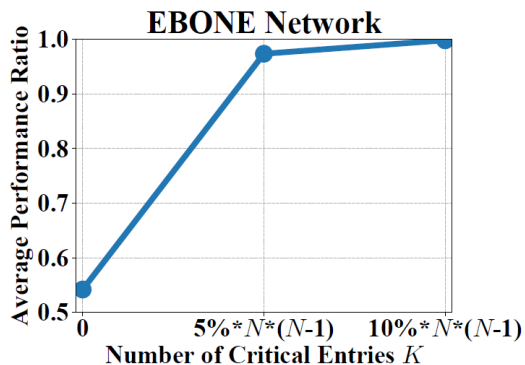

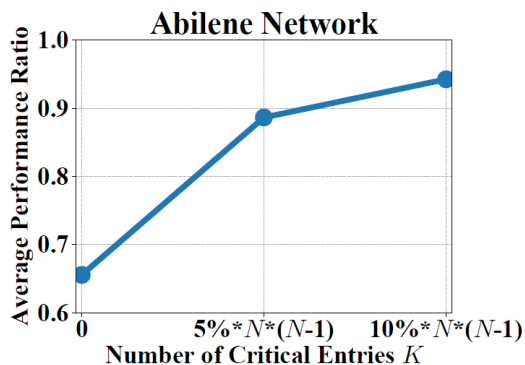
Topology	Nodes	Directed Links
Abilene	12	30
GÉANT	23	74
EBONE	23	76
Tiscali	49	172



$$PR = \frac{18\%}{20\%} = 0.9$$

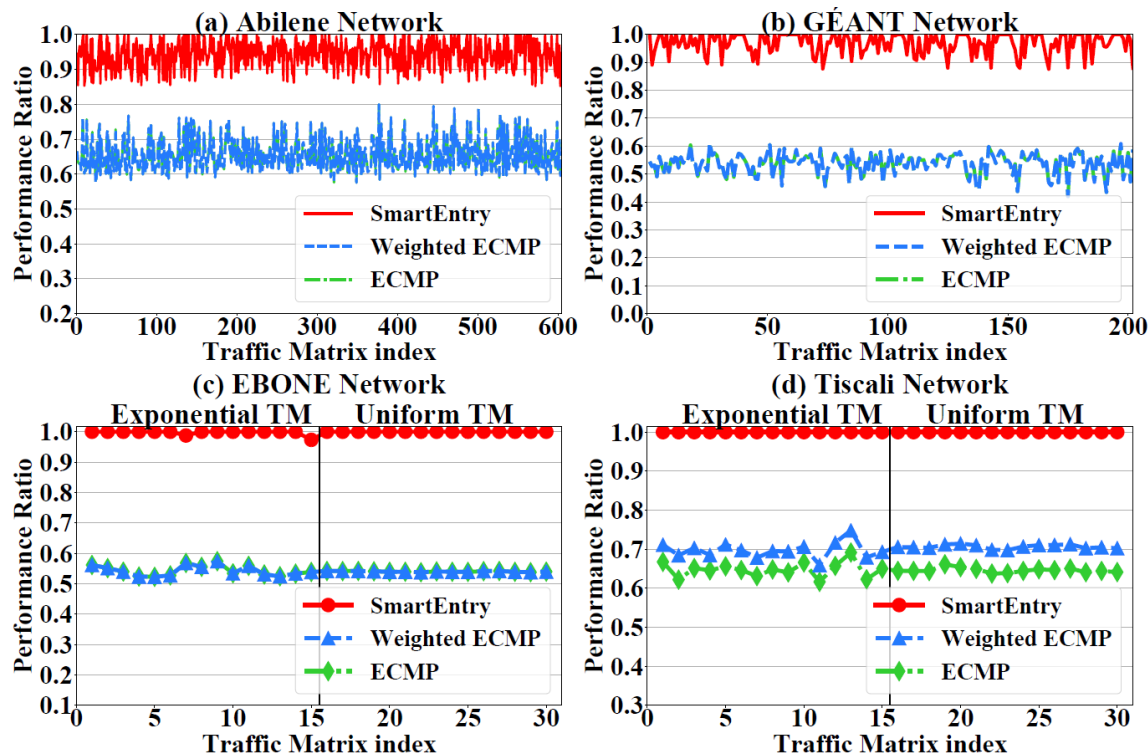
Number of critical entries

- SmartEntry achieves near-optimal performance with only **10% entries** updated



Comparison in different networks

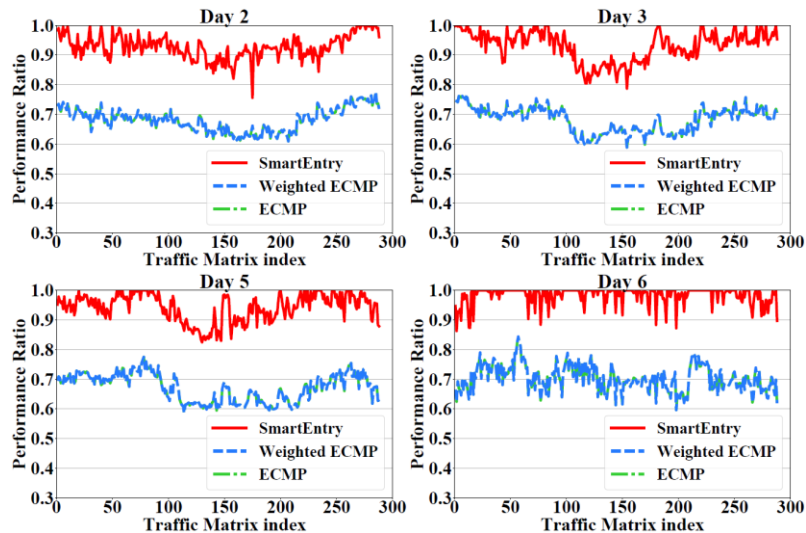
- SmartEntry performs consistently well on real and synthesized traffic matrices




Generalization test

- Training on week 1, test on week 2
- SmartEntry generalizes well to unseen traffics

Abilene Network



GÉANT Network



Conclusion

- With an objective of minimizing maximum link utilization in a network and mitigating routing update overhead, we proposed SmartEntry, a scheme that learns a combination selection policy automatically using reinforcement learning, without any domain specific rule-based heuristic.
- SmartEntry smartly selects a combination of K node-destination pairs for each given traffic matrix and reroutes the selected traffic to achieve load balancing of the network by solving a rerouting optimization problem.
- Extensive evaluations show that SmartEntry achieves near-optimal performance and generalizes well to traffic matrices for which it was not explicitly trained.